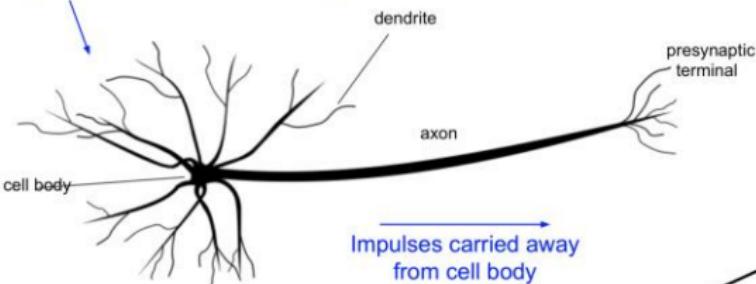


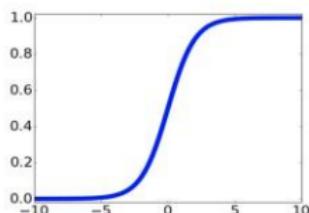
# Neural Network Layers

# Neural network analogy

Impulses carried toward cell body

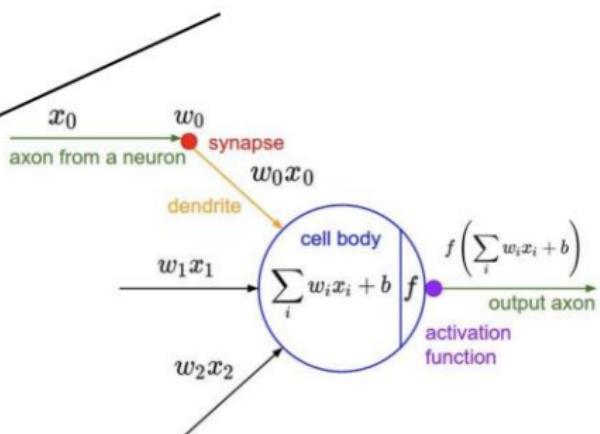


Impulses carried away from cell body

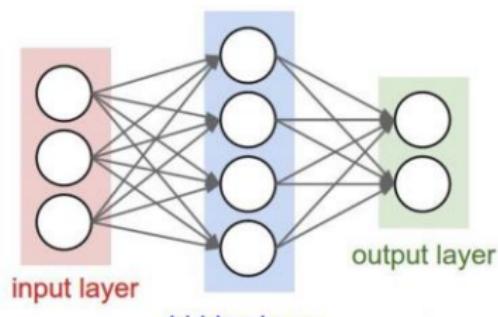


sigmoid activation function

$$\frac{1}{1 + e^{-x}}$$

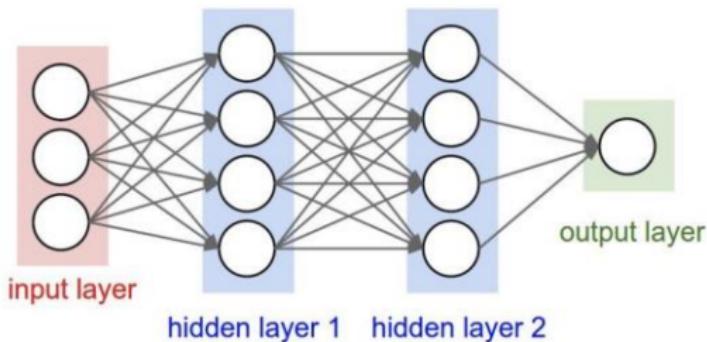


# Neural Network Architectures



"2-layer Neural Net", or  
"1-hidden-layer Neural Net"

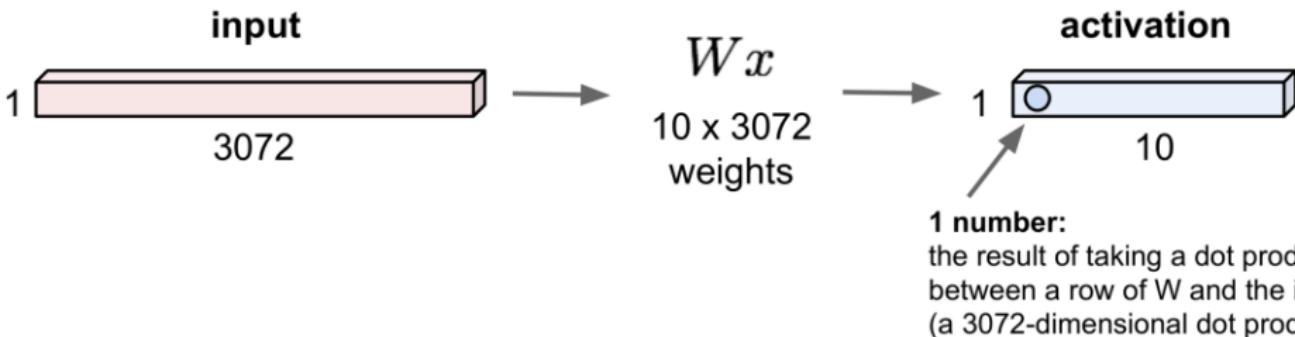
**"Fully-connected" layers**



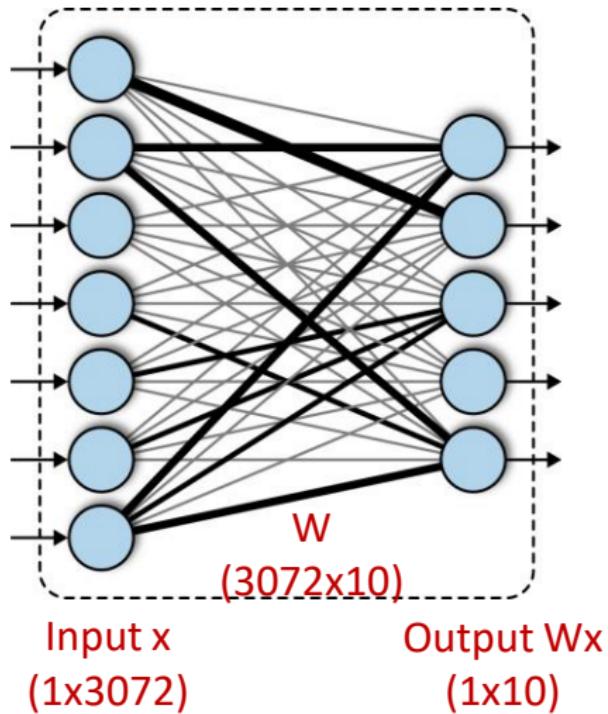
"3-layer Neural Net", or  
"2-hidden-layer Neural Net"

## Fully Connected Layer

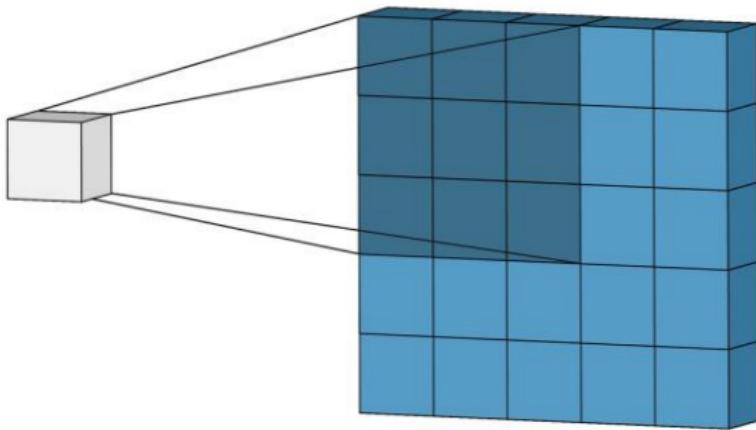
32x32x3 image -> stretch to 3072 x 1



## Fully Connected Layer



## Convolutional layer



# Convolutional layer

$$\begin{matrix} 2 & 4 & 9 & 1 & 4 \\ 2 & 1 & 4 & 4 & 6 \\ 1 & 1 & 2 & 9 & 2 \\ 7 & 3 & 5 & 1 & 3 \\ 2 & 3 & 4 & 8 & 5 \end{matrix} \times \begin{matrix} 1 & 2 & 3 \\ -4 & 7 & 4 \\ 2 & -5 & 1 \end{matrix} = \begin{matrix} 51 & & \\ & & \\ & & \end{matrix}$$

Image                          Filter / Kernel                          Feature

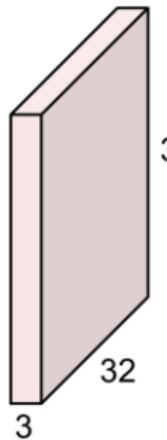
$$\begin{matrix} 2 & 4 & 9 & 1 & 4 \\ 2 & 1 & 4 & 4 & 6 \\ 1 & 1 & 2 & 9 & 2 \\ 7 & 3 & 5 & 1 & 3 \\ 2 & 3 & 4 & 8 & 5 \end{matrix} \times \begin{matrix} 1 & 2 & 3 \\ -4 & 7 & 4 \\ 2 & -5 & 1 \end{matrix} = \begin{matrix} 51 & 66 & 20 \\ 31 & 49 & 101 \\ 15 & 53 & -2 \end{matrix}$$

Image                          Filter / Kernel                          Feature

## Convolutional layer

# Convolution Layer

32x32x3 image



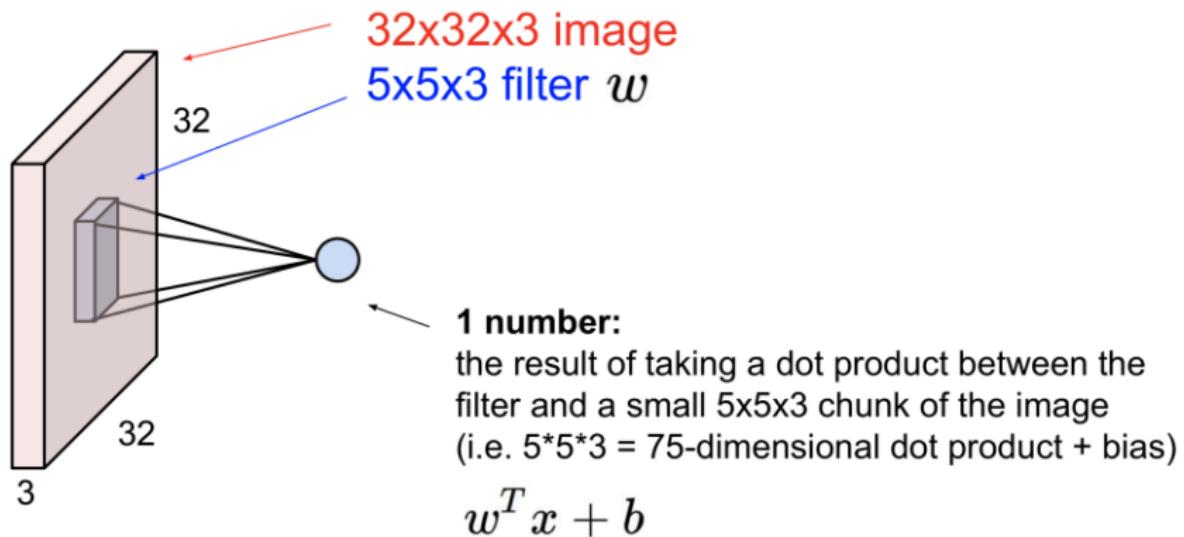
5x5x3 filter



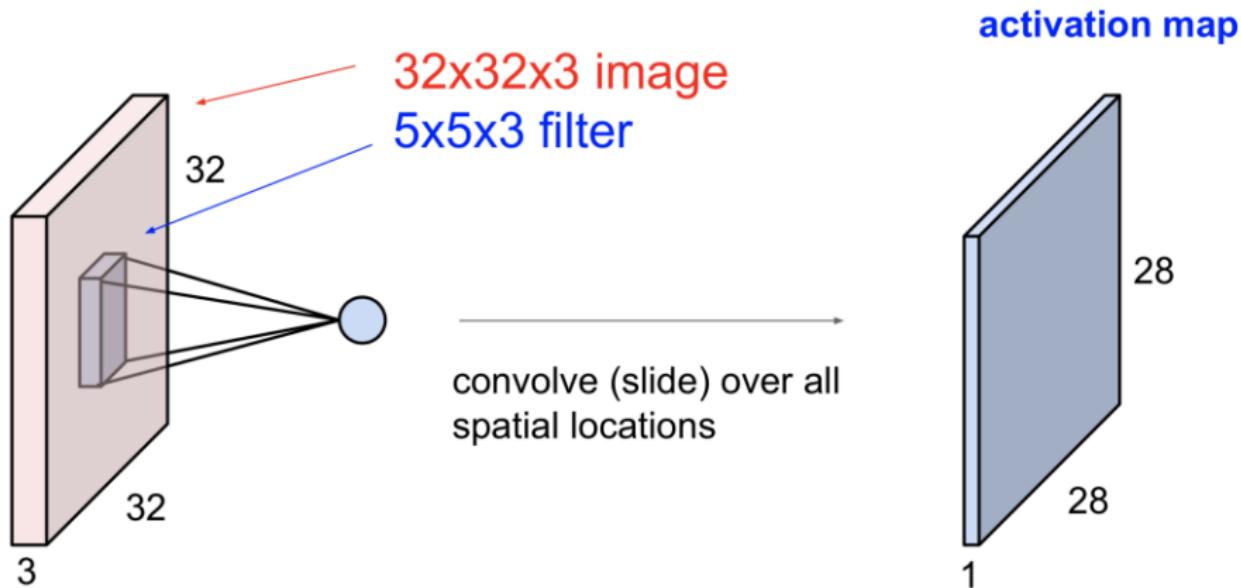
Filters always extend the full depth of the input volume

**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

## Convolutional layer



## Convolutional layer



## Convolutional layer

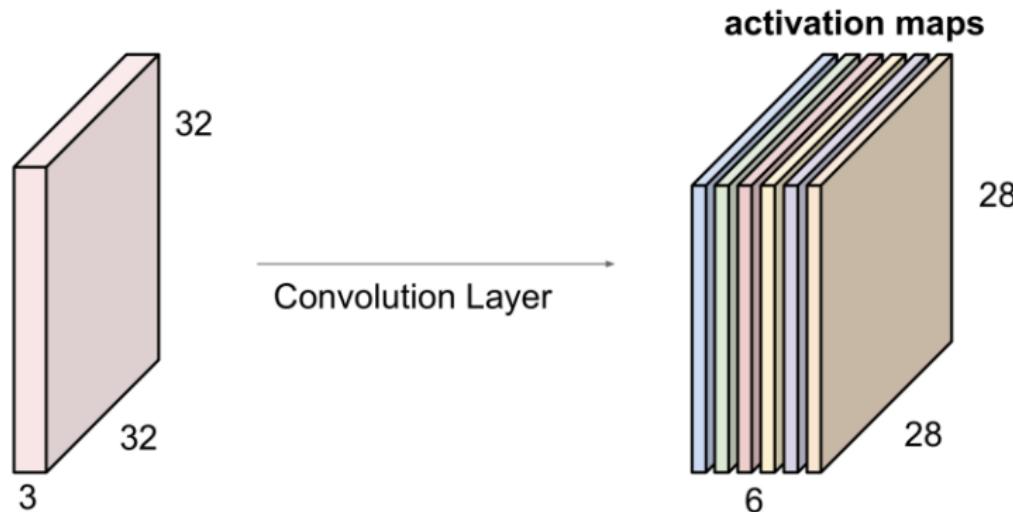
### Convolution Layer

consider a second, green filter



## Convolutional layer

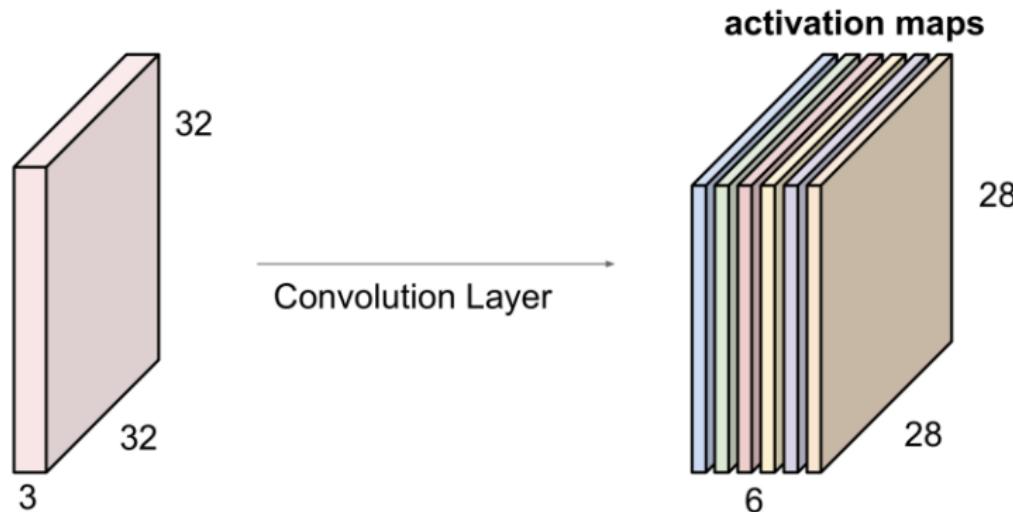
For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size  $28 \times 28 \times 6$ !

## Convolutional layer

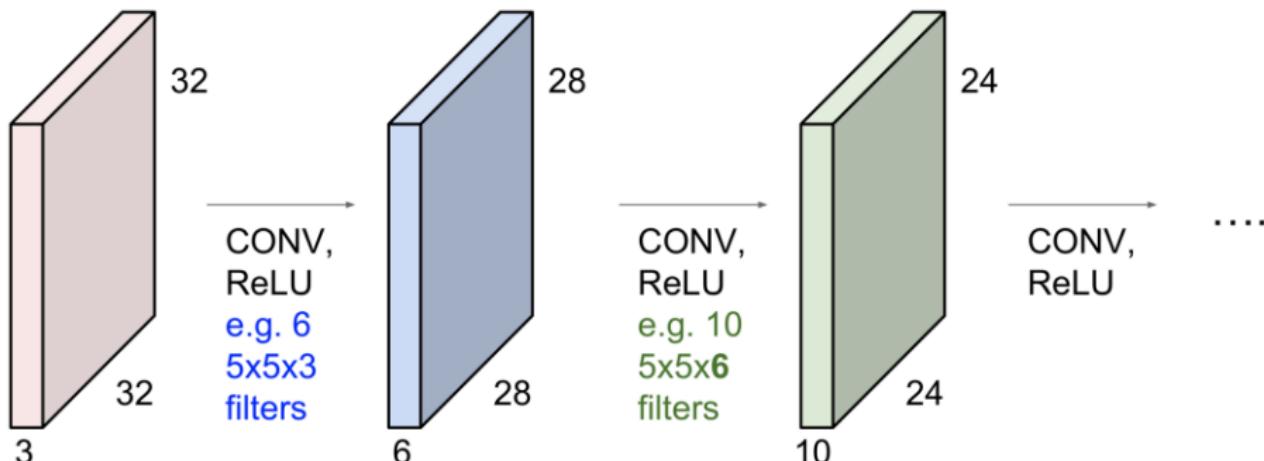
For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size  $28 \times 28 \times 6$ !

## Convolutional layer

**Preview:** ConvNet is a sequence of Convolutional Layers, interspersed with activation functions

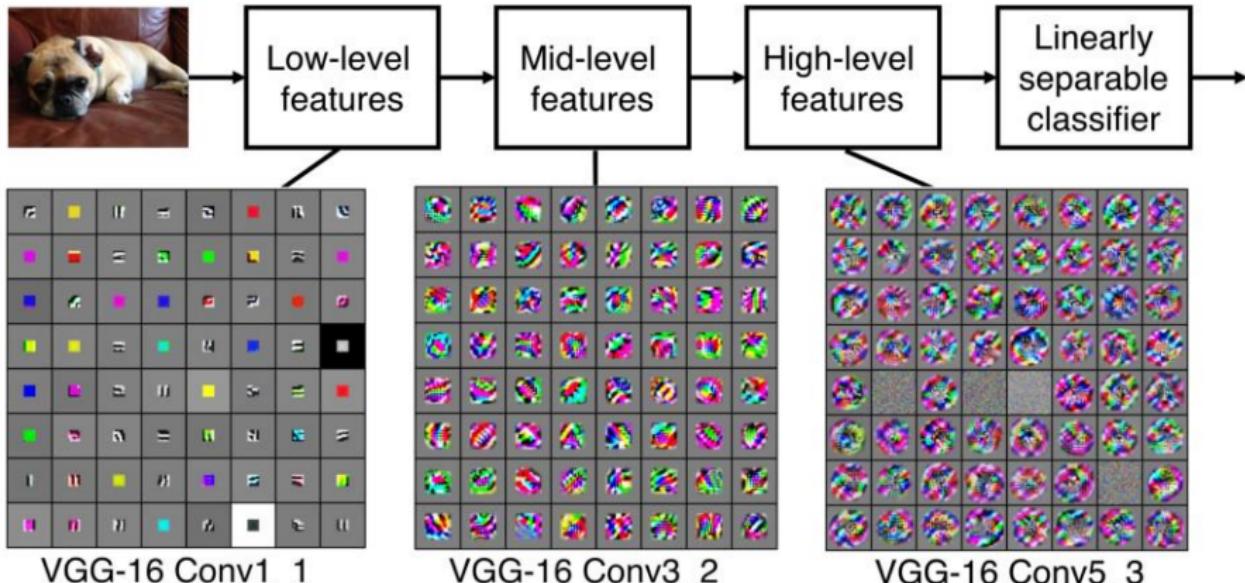


# Convolutional layer

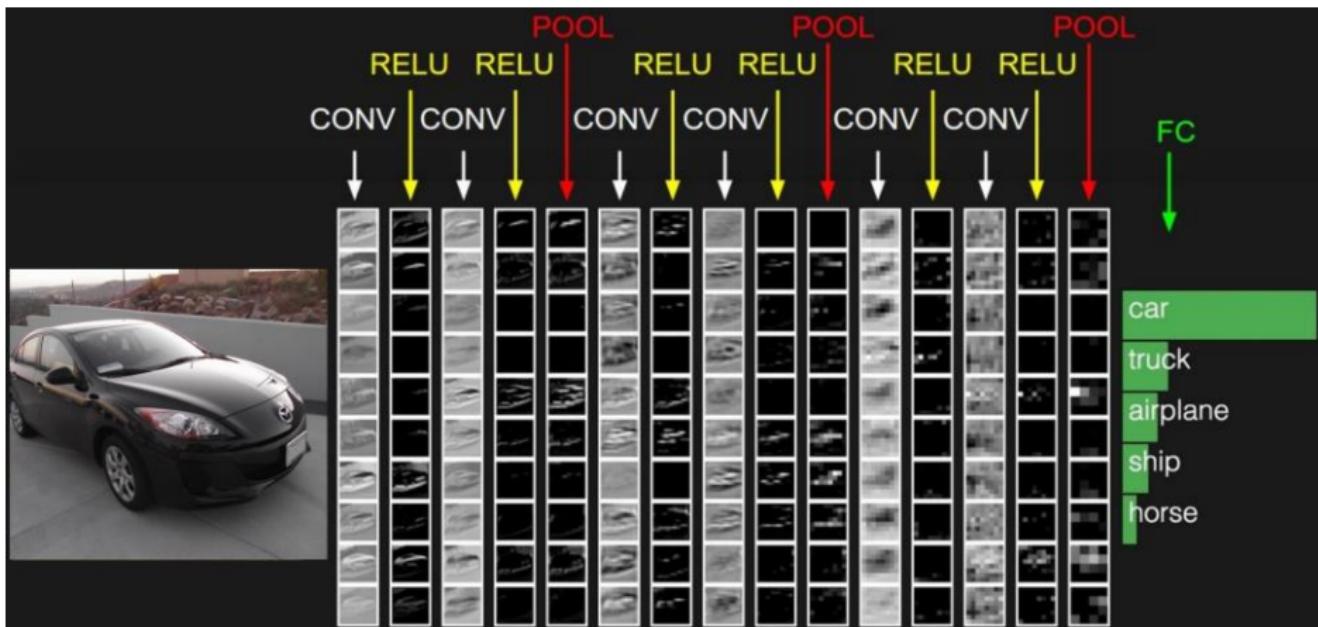
## Preview

[Zeiler and Fergus 2013]

Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].



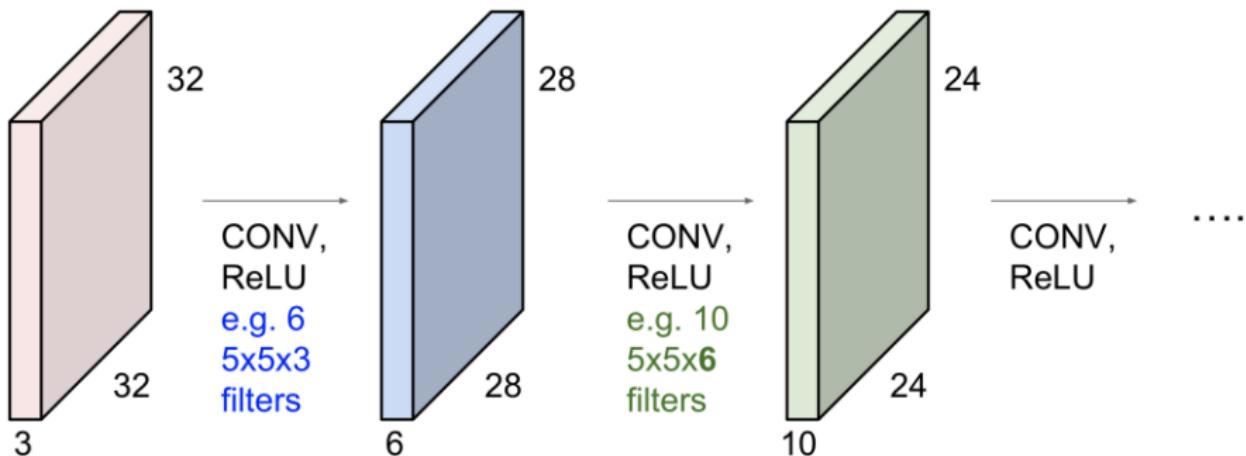
# Convolutional layer



# Convolutional layer

Remember back to...

E.g. 32x32 input convolved repeatedly with 5x5 filters shrinks volumes spatially!  
(32 -> 28 -> 24 ...). Shrinking too fast is not good, doesn't work well.



# Convolutional layer

$$\begin{matrix} \begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 4 & 9 & 1 & 4 & 0 \\ 0 & 2 & 1 & 4 & 4 & 6 & 0 \\ 0 & 1 & 1 & 2 & 9 & 2 & 0 \\ 0 & 7 & 3 & 5 & 1 & 3 & 0 \\ 0 & 2 & 3 & 4 & 8 & 5 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{matrix} & \times & \begin{matrix} 1 & 2 & 3 \\ -4 & 7 & 4 \\ 2 & -5 & 1 \end{matrix} & = & \begin{matrix} 21 & 59 & 37 & -19 & 2 \\ 30 & 51 & 66 & 20 & 43 \\ -14 & 31 & 49 & 101 & -19 \\ 59 & 15 & 53 & -2 & 21 \\ 49 & 57 & 64 & 76 & 10 \end{matrix} \\ \text{Image} & & \text{Filter / Kernel} & & \text{Feature} \end{matrix}$$

- Input 5x5 with padding 1.
- Filter with kernel size 3x3 and stride 1.  
→ Output is 5x5.

## Convolutional layer

Input Volume (+pad 1) (7x7x3)

x[:, :, 0]	0 0 0 0 0 0 0	0 0 0 2 2 2 0	0 2 2 2 2 0 0
0 0 0 0 0 0 0	0 0 0 2 2 2 0	-1 -1 1	= 0
0 0 0 1 1 0 0	0 0 0 2 1 0 1	0 0 1	w0[:, :, 1]
0 0 0 2 2 1 0	0 0 0 2 2 1 0	1 0 -1	= -3
0 0 0 0 0 0 0	0 0 0 0 0 0 0	1 0 0	w0[:, :, 2]
x[:, :, 1]	0 0 0 0 0 0 0	1 -1 0	= 1
0 1 2 0 1 0 0	0 1 0 1 2 2 0	1 -1 1	b0[:, :, 0]
0 1 0 1 2 2 0	0 0 2 2 0 0 0	0 0 0	= 1
0 0 1 1 0 1 0	0 0 0 0 0 0 0	1	Bias b0 (1x1x1)
0 1 1 0 0 0 0	0 0 0 0 0 0 0		b0[:, :, 0]
x[:, :, 2]	0 0 0 0 0 0 0		= 1
0 0 0 2 1 0 1	0 0 2 1 2 2 0		
0 0 2 1 2 2 0	0 2 2 2 2 2 0		
0 0 0 2 0 1 0	0 0 0 2 0 1 0		
0 2 1 2 2 0 0	0 0 0 0 0 0 0		
0 0 0 0 0 0 0			

Filter W0 (3x3x3)

w0[:, :, 0]	0 1 -1
0 1 -1	-1 -1 1
-1 0 1	

Filter W1 (3x3x3)

w1[:, :, 0]	0 1 -1
0 -1 1	0 1 0
0 1 0	w1[:, :, 1]
0 1 -1	1 0 -1
1 0 -1	1 -1 1
1 -1 1	

Output Volume (3x3x2)

o[:, :, 0]	3 -1 -5
1 3 0	-2 0 0
-2 0 0	o[:, :, 1]
-1 -2 -4	4 3 3
4 3 3	-1 8 0
-1 8 0	

Bias b0 (1x1x1)

b0[:, :, 0]	0
-------------	---

toggle movement

Visualization of how a convolutional layer works  
available at <https://cs231n.github.io/convolutional-networks/>

Visualization of how a convolutional layer works. Animation available at <https://cs231n.github.io/convolutional-networks/>

19

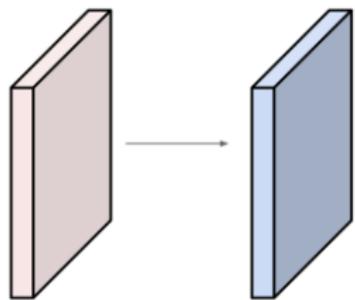
## Convolutional layer

Examples time:

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2

Output volume size: ?

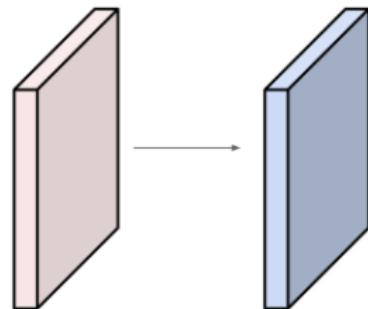


## Convolutional layer

Examples time:

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2



Output volume size:

$(32+2*2-5)/1+1 = 32$  spatially, so

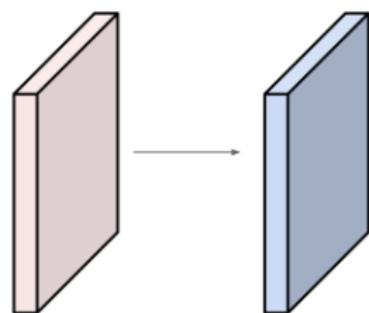
**32x32x10**

## Convolutional layer

Examples time:

Input volume: **32x32x3**

10 **5x5** filters with stride 1, pad 2



Number of parameters in this layer?

each filter has  $5 \times 5 \times 3 + 1 = 76$  params      (+1 for bias)  
=> **76 \* 10 = 760**

# Convolutional layer

Summary of the Convolutional layer:

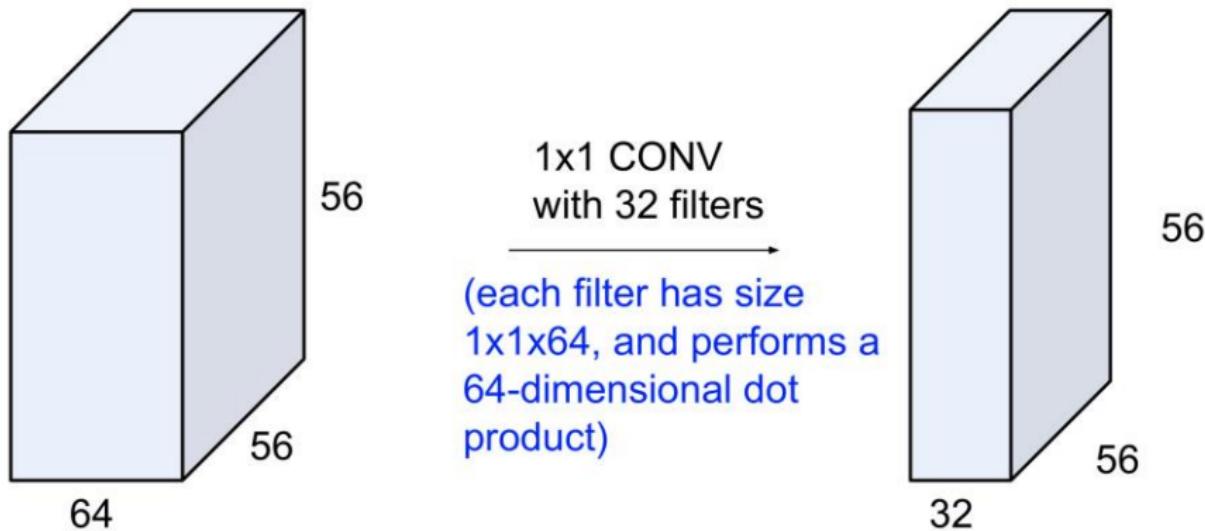
- Accept a volume of size  $W_1 \times H_1 \times D_1$
- Require four hyperparameters:
  - Spatial extent  $F$
  - Number of filters  $K$
  - Stride  $S$
  - Zero padding  $P$
- Produce a volume of size  $W_2 \times H_2 \times D_2$  where:
  - $W_2 = (W_1 + 2P - F)/S + 1$
  - $H_2 = (H_1 + 2P - F)/S + 1$
  - $D_2 = K$
- With parameter sharing, it introduces  $F \times F \times D_1$  weights per filter, for a total of  $(F \times F \times D_1) \times K$  weights and  $K$  biases.
- In the output volume, the  $d$ -th depth slice (of size  $W_2 \times H_2$ ) is the result performing a valid convolution of the  $d$ -th filter over the input volume with a stride  $S$ , and then offset by  $d$ -th bias.

## Convolutional layer

**Question:** What is the reason for which the Convolutional layer was born? Does it make sense to have a Convolutional layer with kernel size 1x1? Why or why not?

## One-by-one convolution layer

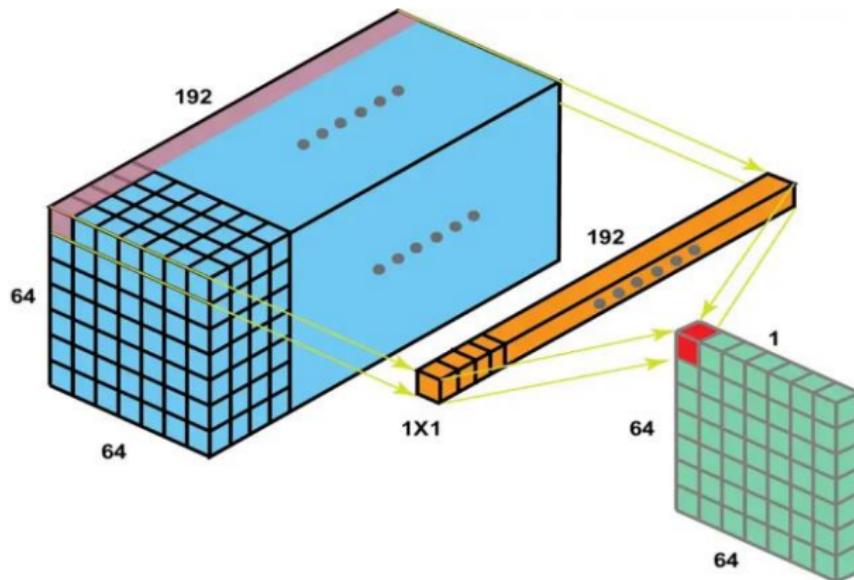
(btw, 1x1 convolution layers make perfect sense)



## One-by-one Convolutional layer

One-by-one (1x1) convolutional layer is used for:

- Dimensionality reduction.
- Dimensionality augmentation.



## One-by-one Convolutional layer

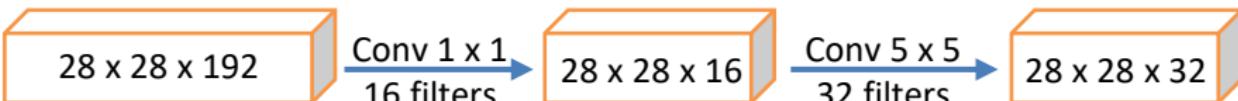
One-by-one (1x1) convolutional layer is used for:

- **Dimensionality reduction.**
- Dimensionality augmentation.



Number of learnable parameters = ?  
Number of operations = ?

Assume that there is **no bias** for the sake of simplicity.



Number of learnable parameters = ?  
Number of operations = ?

Assume that there is **no bias** for the sake of simplicity.

## One-by-one Convolutional layer

One-by-one (1x1) convolutional layer is used for:

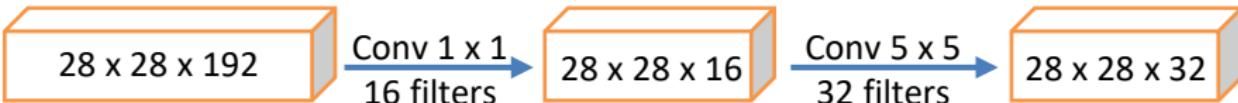
- **Dimensionality reduction.**
- Dimensionality augmentation.



Number of learnable parameters =  $(5 \times 5 \times 192) \times 32 = 153,600$  params.

Number of operations =  $(5 \times 5 \times 192) \times (28 \times 28) \times 32 = 120,422,400$  ops.

---



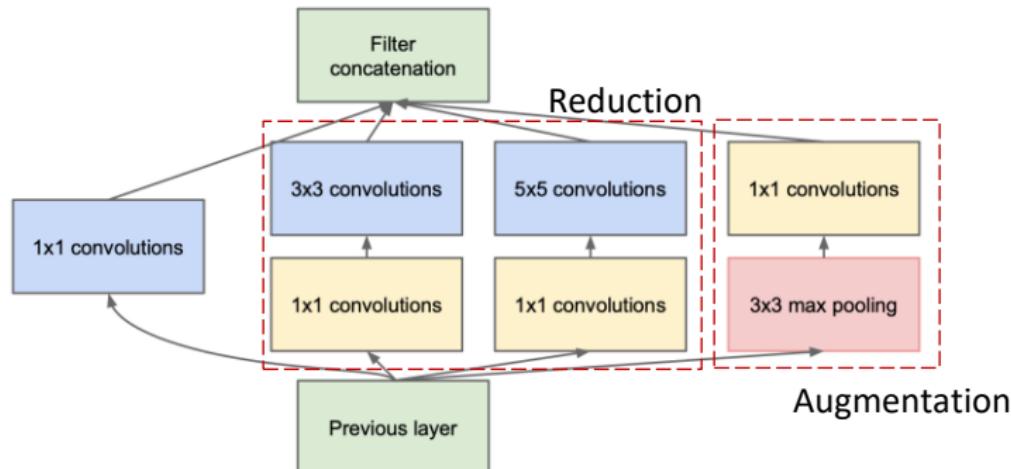
Number of learnable parameters =  $(1 \times 1 \times 192) \times 16 + (5 \times 5 \times 16) \times 32 = 15,872$ .

Number of operations =  $(1 \times 1 \times 192) \times (28 \times 28) \times 16 + (5 \times 5 \times 16) \times (28 \times 28) \times 32$   
= 12,443,648 ops.

## One-by-one Convolutional layer

One-by-one (1x1) convolutional layer is used for:

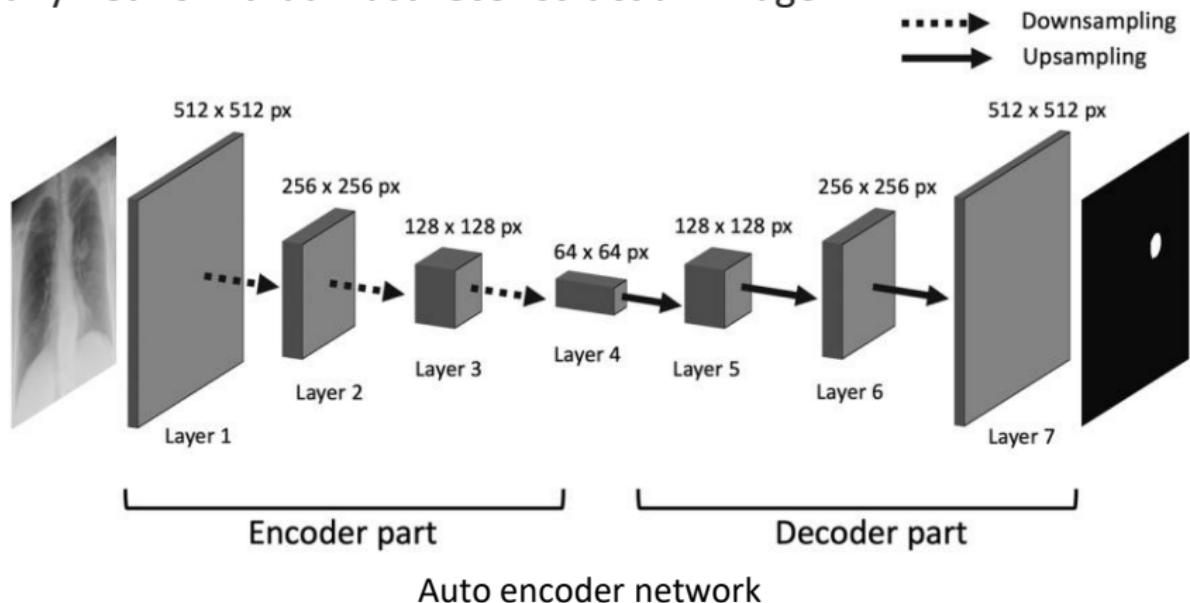
- Dimensionality reduction (more **popular** than augmentation).
- **Dimensionality augmentation:** increase the number of feature maps after pooling, artificially creating *more projections* of the *down-sampled* feature map content.



**Inception** module with 1x1 conv for dimensionality reduction and augmentation.

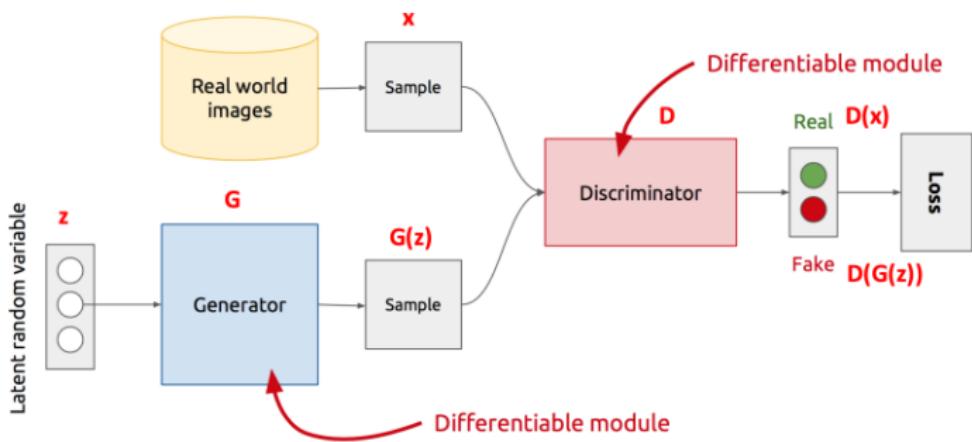
## Transposed convolutional layer

- Convolutional layer: **down-samples** the input.
- Transposed convolutional layer: **up-samples** the spatial dimensions, usually used in **auto-encoders** and **GANs**, or generally any network that must **reconstruct** an image.



# Transposed convolutional layer

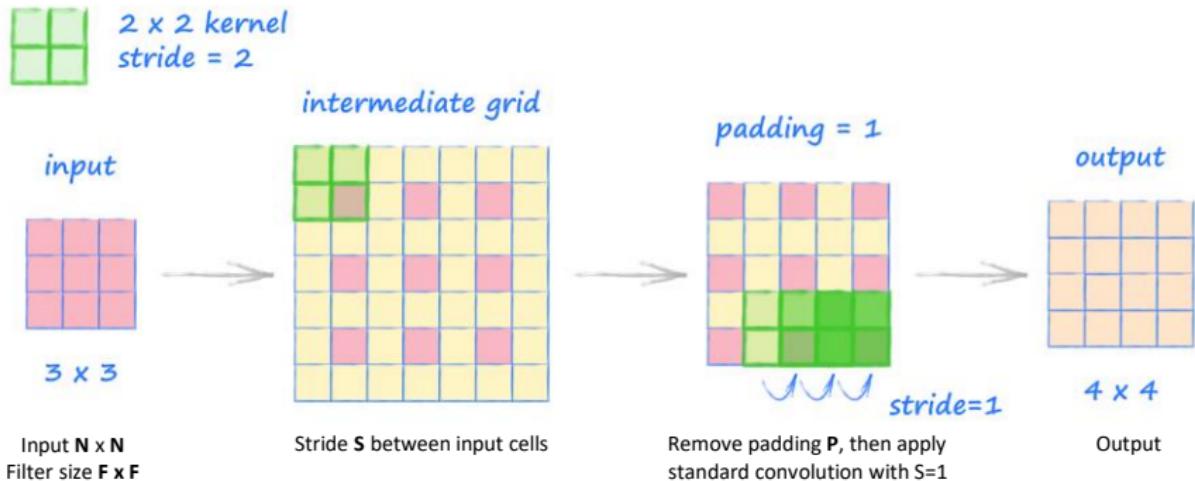
- Convolutional layer: **down-samples** the input.
- Transposed convolutional layer: **up-samples** the spatial dimensions, usually used in **auto-encoders** and **GANs**, or generally any network that must **reconstruct** an image.



Generative Adversarial Network (GAN)

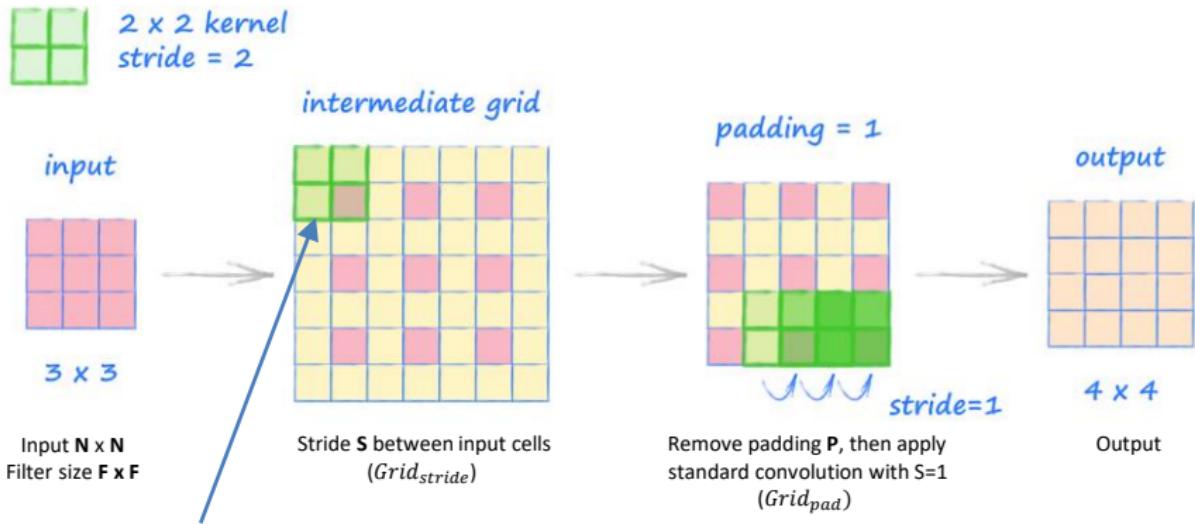
Transposed convolutional layers are used in the Generator (G).

# Transposed convolutional layer



**Question:** what is the output size of a transposed convolutional layer,  
given input  $N \times N$ , stride  $S$  and padding  $P$ ?

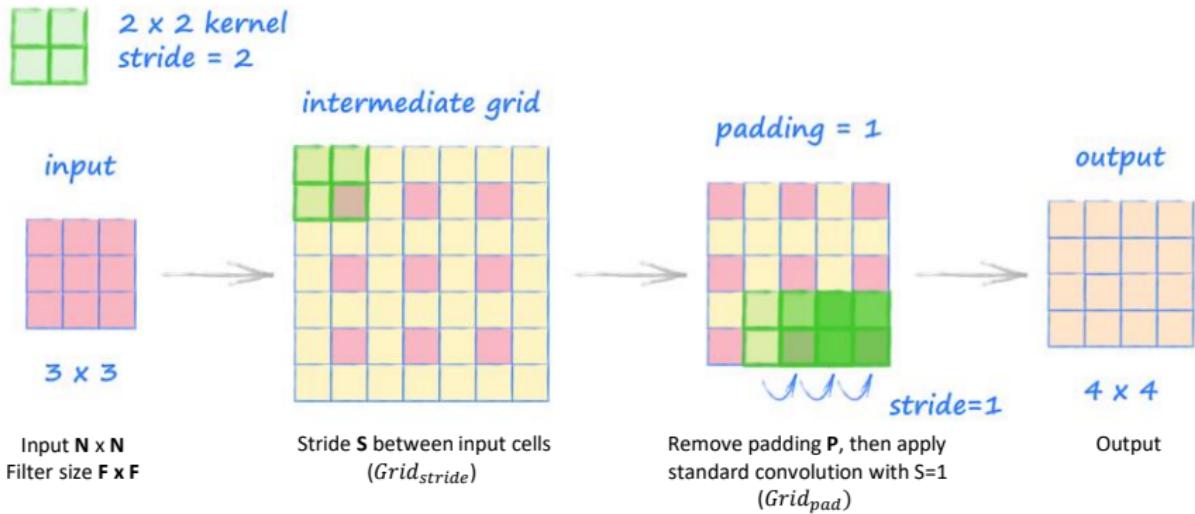
# Transposed convolutional layer



$$Grid_{stride} = N + (N - 1) \times (S - 1) + (F - 1) \times 2$$

$$\begin{aligned} Grid_{pad} &= Grid_{stride} - 2 \times P = N + (N - 1) \times (S - 1) + (F - 1) \times 2 - 2 \times P \\ &= (N - 1)S + 2F - 2P - 1 \end{aligned}$$

# Transposed convolutional layer

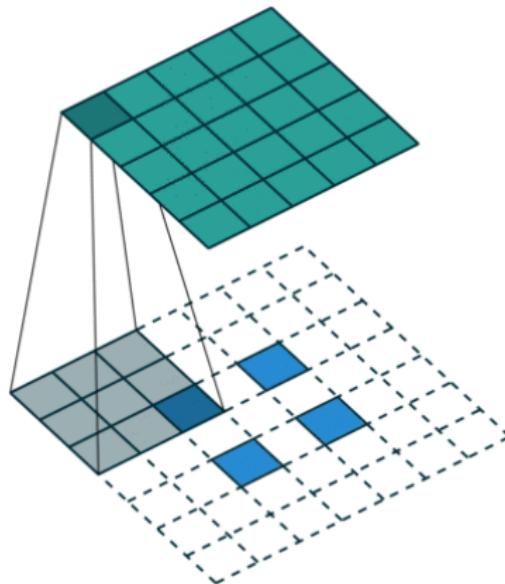


$$Output = \frac{Grid_{pad} + 2 \times P' - F}{S'} + 1 = Grid_{pad} - F + 1 \text{ where } P' \text{ is always } 0 \text{ and } S' \text{ is always } 1.$$

$$\textbf{Output} = (N - 1)S + 2F - 2P - 1 - F + 1 = (\textbf{N} - \textbf{1})\textbf{S} + \textbf{F} - \textbf{2P}.$$

Above example: Output =  $(3 - 1)2 + 2 - 2 = 4$ . Hence, the output feature map is  $4 \times 4$ .

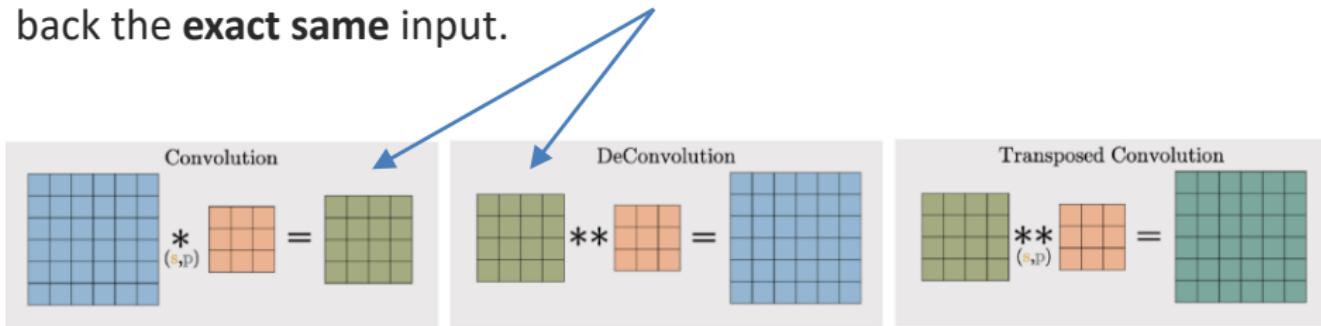
## Transposed convolutional layer



**Question:** In this transposed convolutional layer,  
what are input size **N**, stride **S**, filter size **F** and padding **P**?

# Deconvolution

**Deconvolution** is a mathematical operation that reverses the process of a convolutional layer, i.e., get an input through a convolutional layer, then *get its output through a deconvolutional layer*, we get back the **exact same** input.



Deconvolution is not transposed convolution. People often misinterpret “transposed convolution” as “deconvolution”.

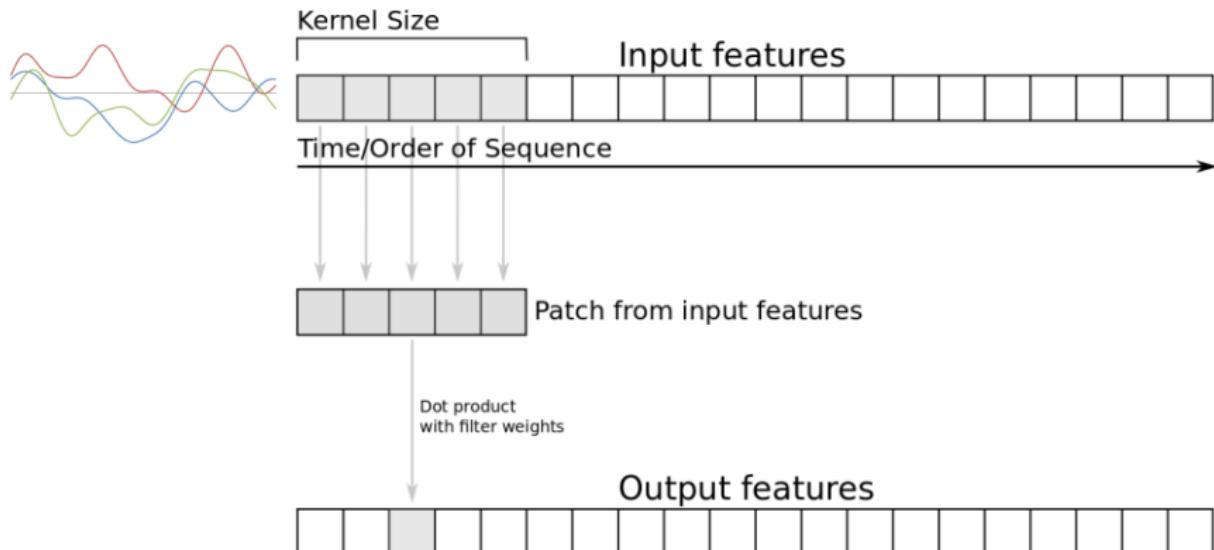
## Dimensions of convolutions

- **1D Convolution:** kernel moves in **1 direction**. Input and output data of 1D CNN is 2 dimensional. Mostly used on time-series data.
- **2D Convolution:** kernel moves in **2 directions**. Input and output data of 2D CNN is 3 dimensional. Mostly used on image data.
- **3D Convolution:** kernel moves in **3 directions**. Input and output data of 3D CNN is 4 dimensional. Mostly used on 3D image data (MRI, CT Scans, Video).

**1D convolution is not 1x1 convolution.**

# 1D Convolution

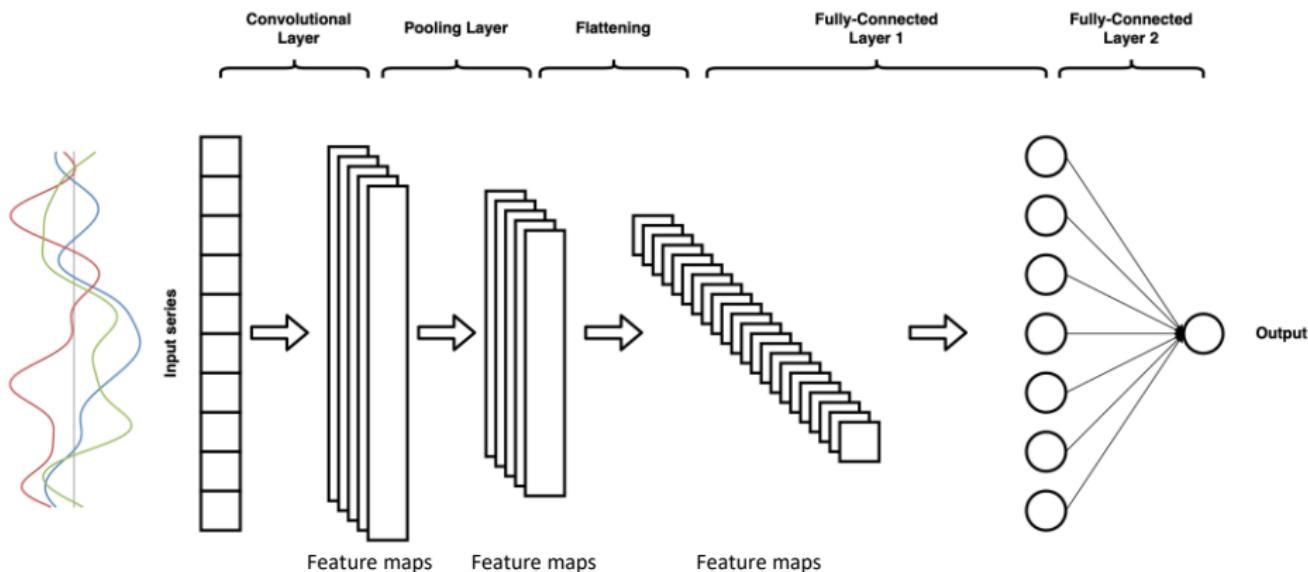
- **1D Convolution:** kernel moves in **1 direction**. Input and output data of 1D CNN is 2 dimensional. Mostly used for **forecasting time-series/sequential data**.



1D convolutional layer operates on time-series/sequential data

# 1D Convolution

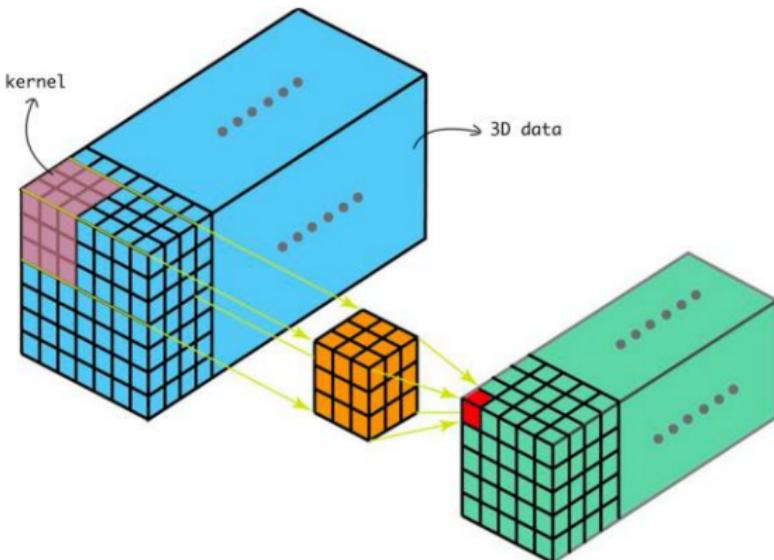
- **1D Convolution:** kernel moves in **1 direction**. Input and output data of 1D CNN is 2 dimensional. Mostly used for **forecasting time-series/sequential data**.



An example of neural network architecture to process time-series/sequential data.

## Three-Dimensional (3D) Convolutional layer

- **1D Convolution:** kernel moves in **1 direction**. Mostly used on **time-Series** data.
- **2D Convolution:** kernel moves in **2 directions**. Mostly used on **image data**.
- **3D Convolution:** kernel moves in **3 directions**. Mostly used on **3D image data** (MRI, CT Scans, Video).

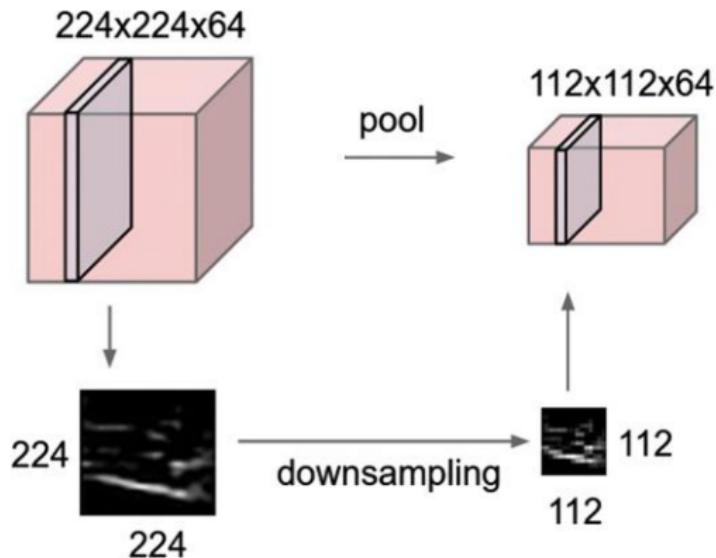


3D convolutional layer operates on 3D image data

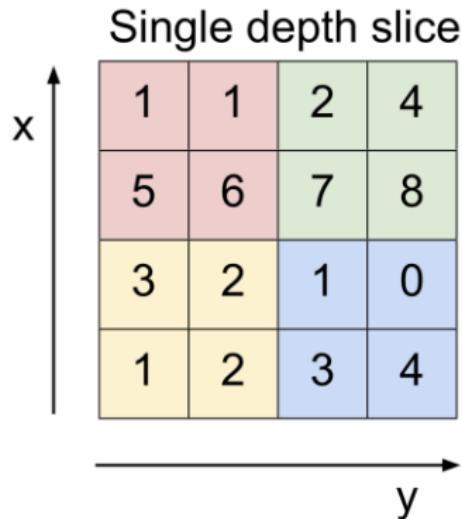
## Pooling layer

Pooling layer:

- makes the representations **smaller** and more manageable.
- operates on each activation map independently.



## Pooling layer



max pool with 2x2 filters  
and stride 2

The output matrix is:

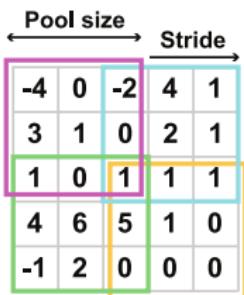
6	8
3	4

Example of max pooling operation

# Pooling layer

Various pooling layers:

- max pooling
- min pooling
- average pooling
- sum pooling
- ...



Max Pooling



Min Pooling



Average Pooling

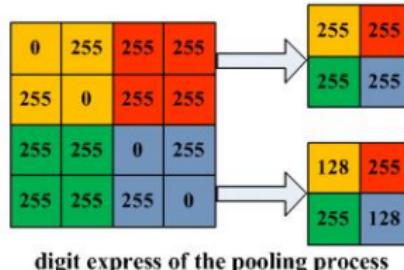
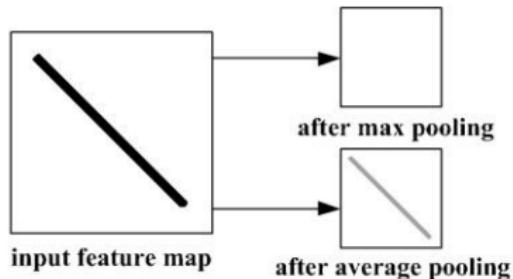


Features

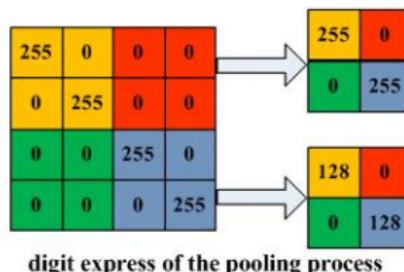
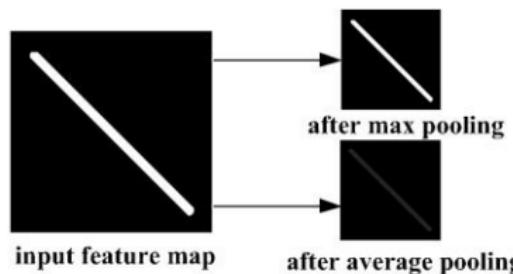
Output

Example of different types of pooling layers.

# Pooling layer



(a) Illustration of max pooling drawback



(b) Illustration of average pooling drawback

Comparison between max pooling and average pooling.

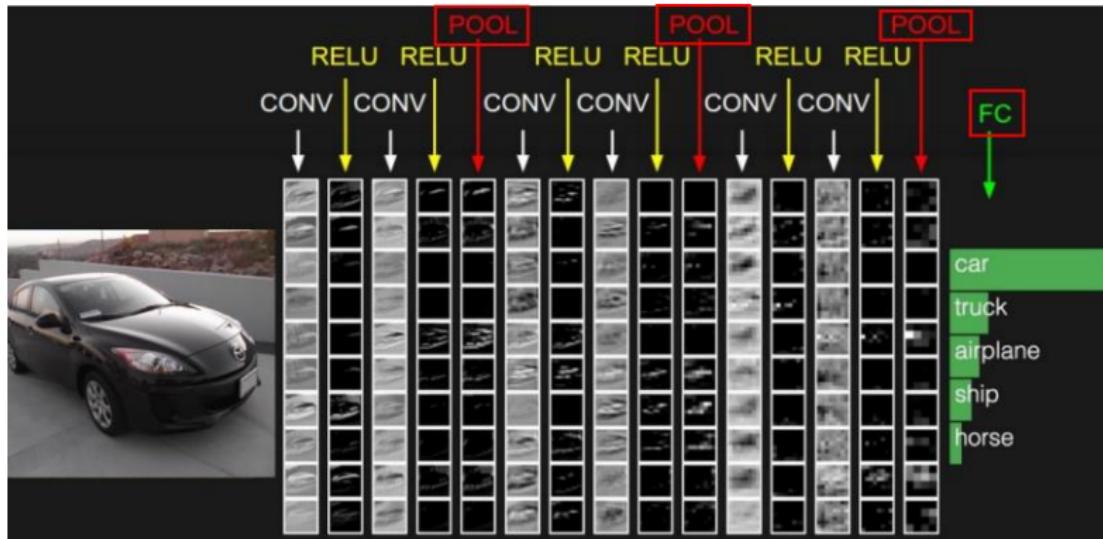
## Pooling layer

Summary of the Pooling layer:

- Accept a volume of size  $W_1 \times H_1 \times D_1$
- Require three hyperparameters:
  - Spatial extent  $F$
  - Stride  $S$
- Produce a volume of size  $W_2 \times H_2 \times D_2$  where:
  - $W_2 = (W_1 - F)/S + 1$
  - $H_2 = (H_1 - F)/S + 1$
  - $D_2 = D_1$
- It introduces zero parameters since it computes a fixed function of the input.
- Note that it is not common to use zero-padding for Pooling layer.

# Fully Connected Layer (FC Layer)

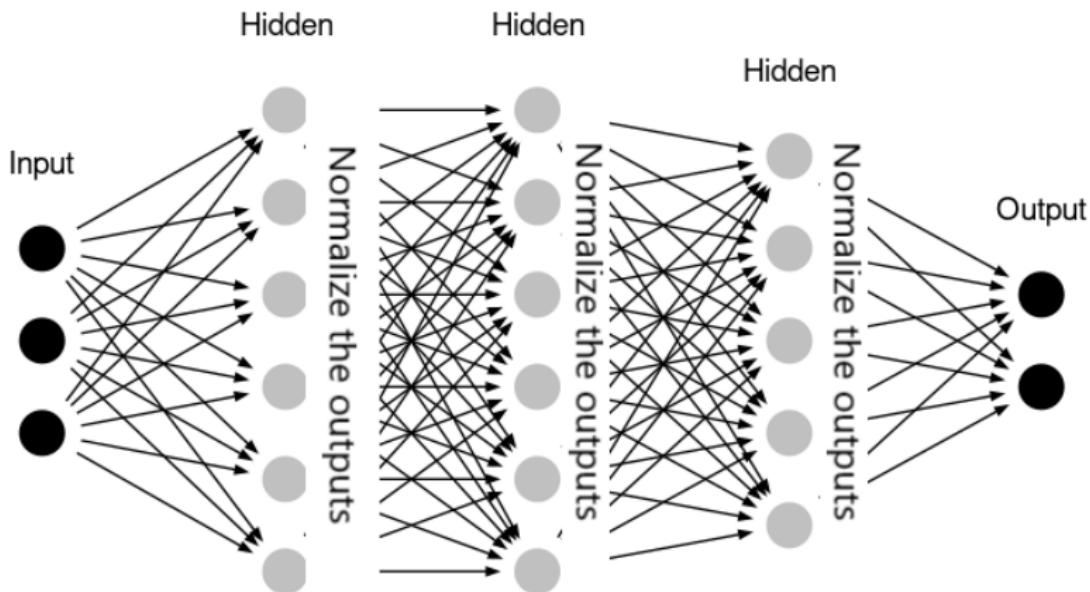
**FC layer:** contains neurons that connect to the entire input volume.



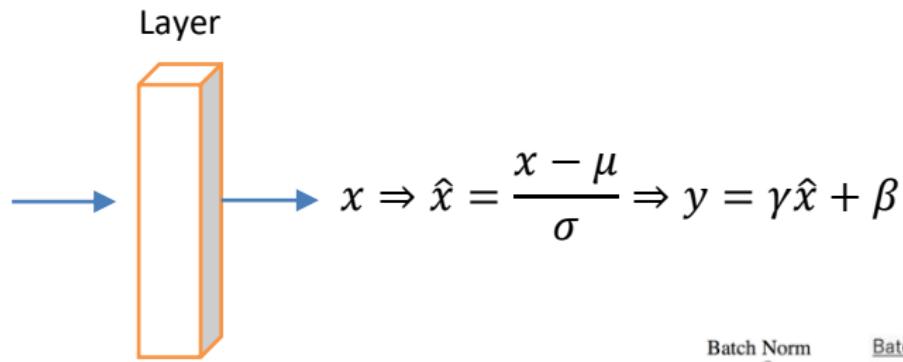
A conventional deep neural architecture for classification problem.

## Batch Normalization

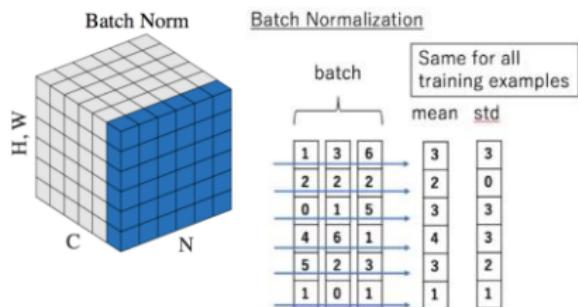
Idea: if the input is normalized, why not at all intermediate layers?



# Batch Normalization



$\mu$ : mean of  $x$  in mini-batch.  
 $\sigma$ : std of  $x$  in mini-batch.  
 $\gamma$ : scale, learnable.  
 $\beta$ : shift, learnable.



## Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots m\}$ ;  
Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

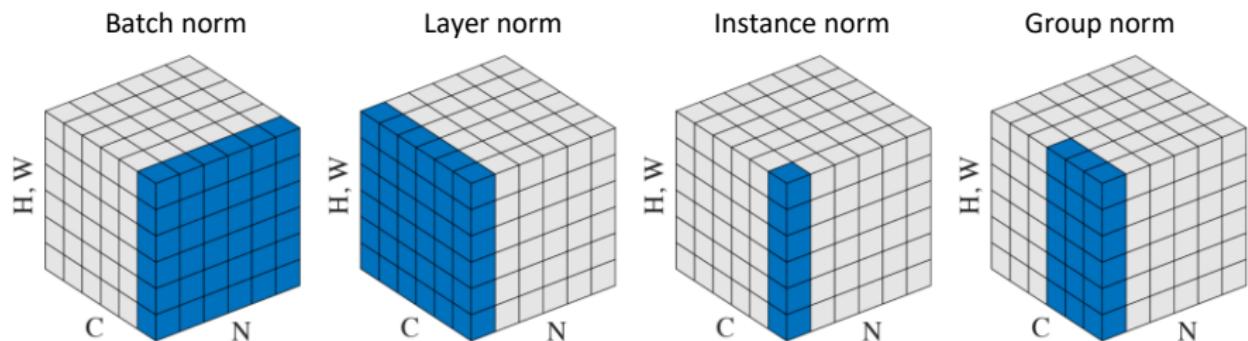
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Normalization methods

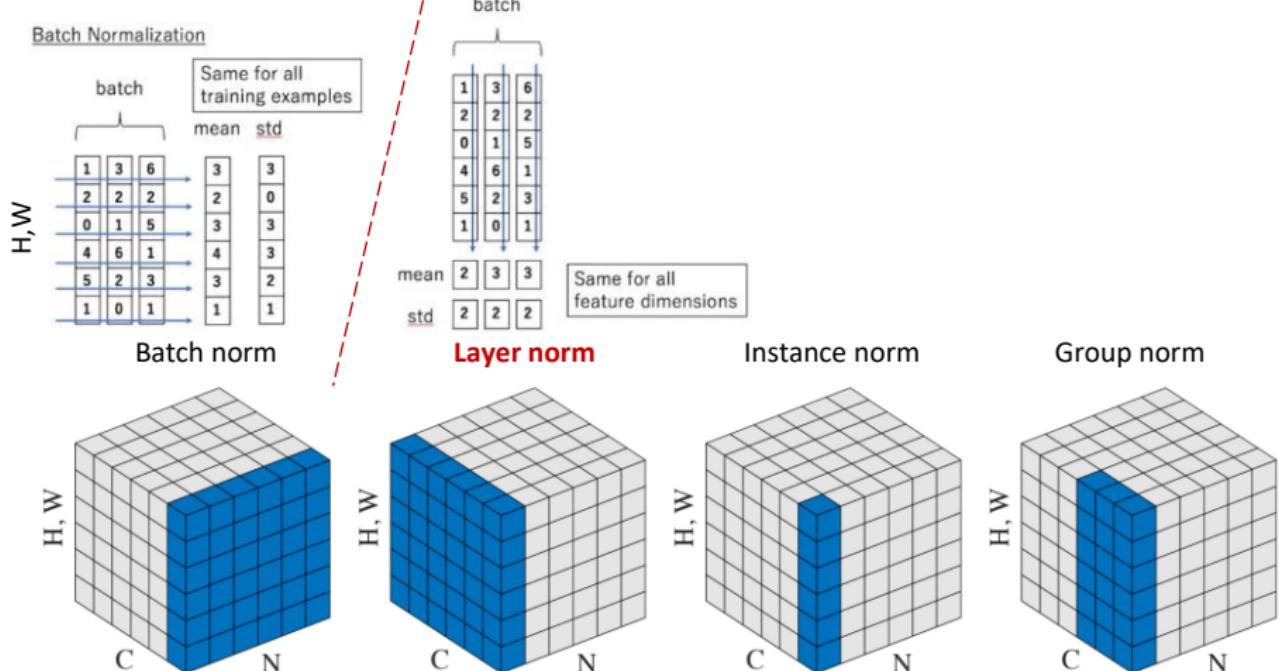
Various normalization layers:

- batch normalization
- layer normalization
- instance normalization
- group normalization



Normalization methods. Each subplot shows a feature map tensor, with **N** as the batch axis, **C** as the channel axis, and **(H, W)** as the spatial axes. The pixels in blue are normalized by the same mean and variance, computed by aggregating the values of these pixels.

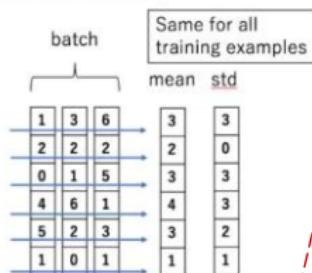
# Normalization methods



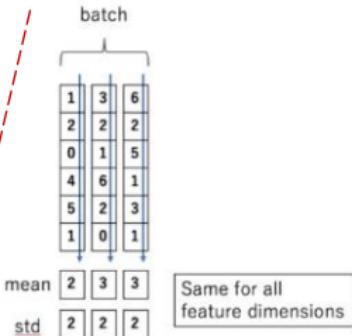
**Layer normalization:** normalizes the activations of the previous layer for each given data sample independently in a batch. This normalization is often used in RNN to stabilize the network convergence.

# Normalization methods

Batch Normalization

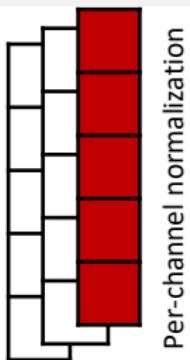


Batch norm



Layer norm

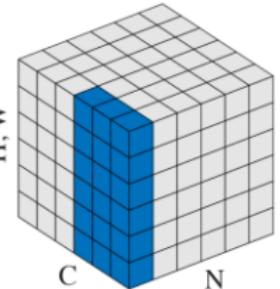
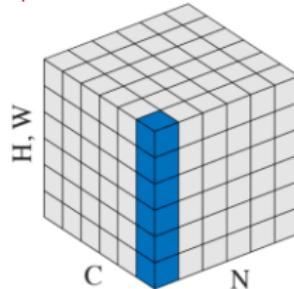
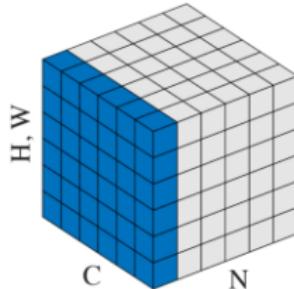
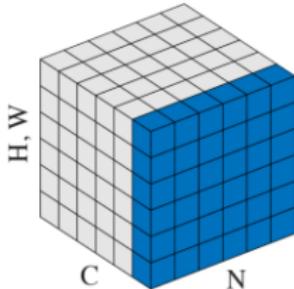
One sample  
Channels



Per-channel normalization

Instance norm

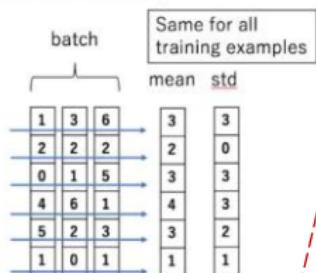
Group norm



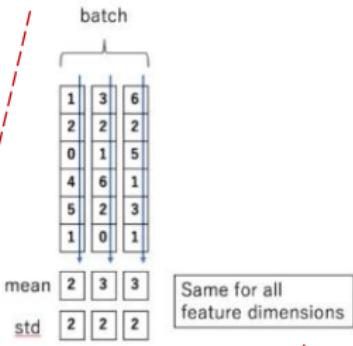
**Instance normalization:** normalizes all features of each channel of a given example independently, to remove image-content instance-specific mean and covariance shift across samples in a data batch.

# Normalization methods

Batch Normalization

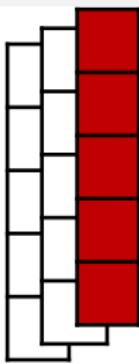


Batch norm



Layer norm

One sample  
Channels

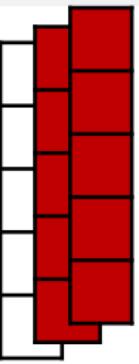


Instance norm

Per-channel normalization

Group norm

One sample  
Channels



H, W

H, W

C N

H, W

C N

H, W

C N

H, W

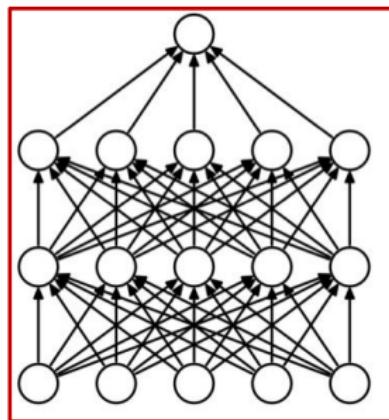
C N

**Group normalization:** divides channels into groups and normalizes the features within each group. Group norm is a generalization of Instance norm where group size = 1. The motivation is features may span across a group of channels.

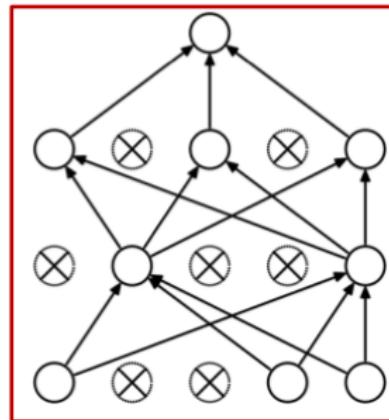
# Dropout

What is **dropout**?

How does **dropout** work during **training** and during **testing**?



Standard neural network

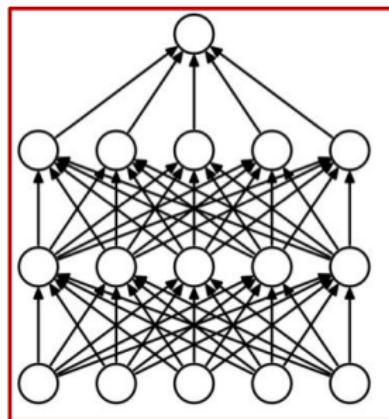


After applying dropout

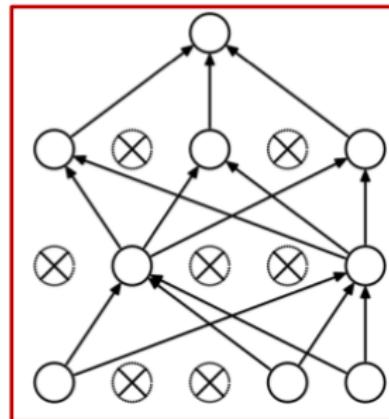
# Dropout

What is dropout?

- **Randomly** “drop” the neurons.
- **Prevent overfitting** by reducing co-adaptation of neurons.
- Like training many **random sub-networks**.



Standard neural network

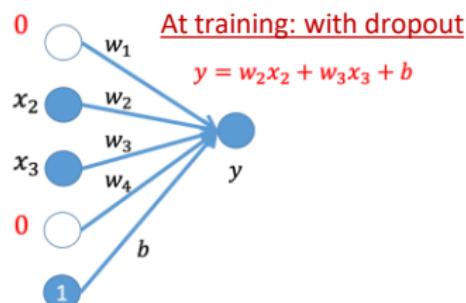
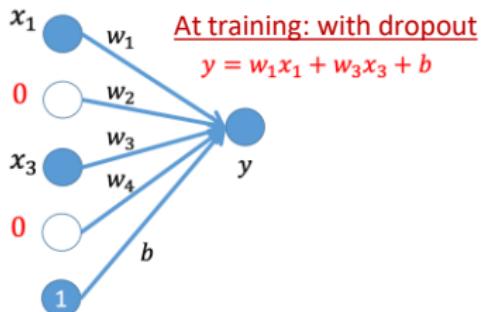
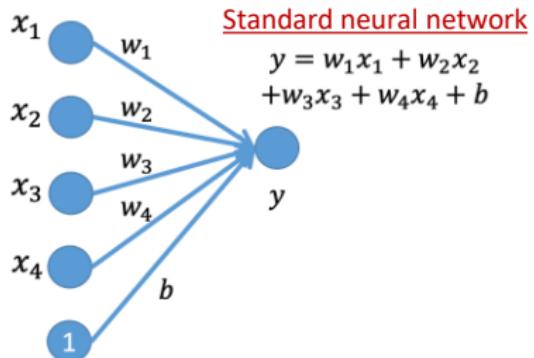


After applying dropout

# Dropout

## Dropout during training:

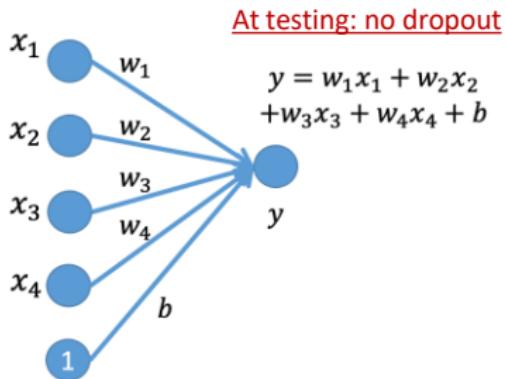
- For each neuron, randomly set its input to zero with probability p, for example p=0.5.



# Dropout

Dropout **during testing**:

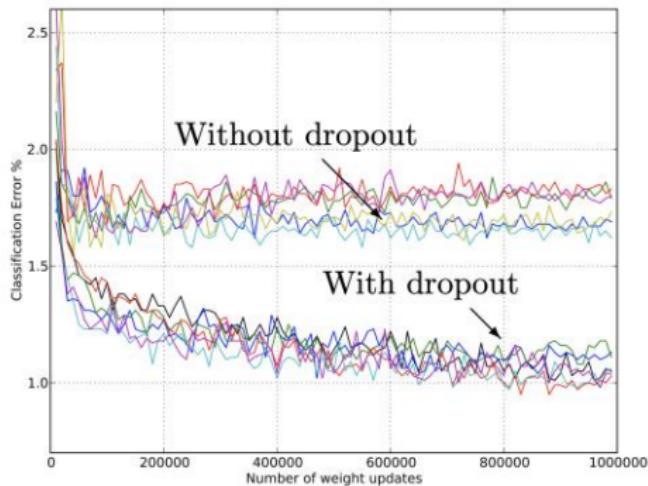
- **No dropout**, i.e., set all neurons with dropout probability p=0.
- Note: Dropouts can be used for neural network inference by dropping during predictions and predicting multiple times to get a distribution.



# Dropout

Efficiency of dropout:

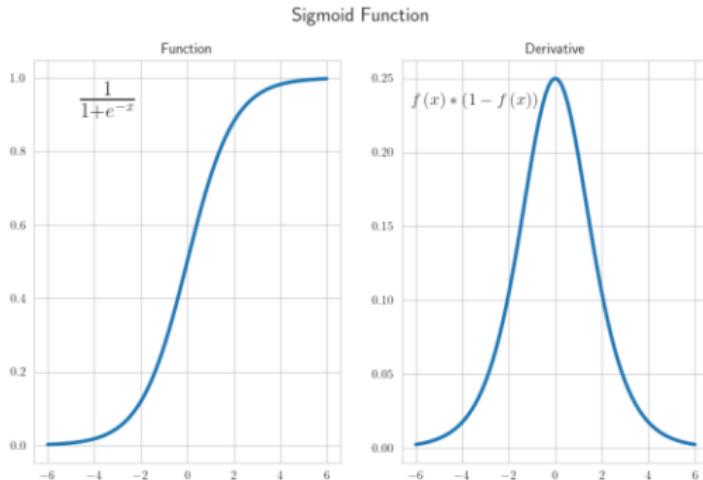
- Widely used and **highly effective**.
- Proposed as an alternative to ensemble methods, which is too expensive for neural networks.



Test error for different architectures with and without dropout.

The networks have 2 to 4 hidden layers each with 1024 to 2048 units.

# Activation functions - Sigmoid



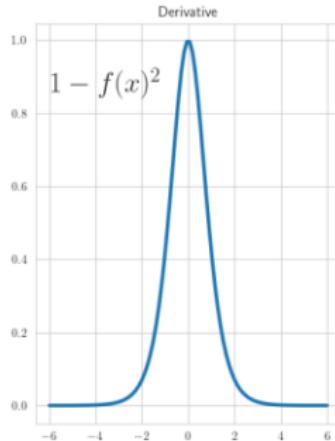
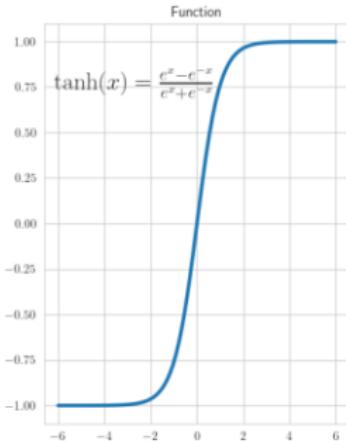
$$f(x) = \frac{1}{1 + e^{-x}}$$

$$f'(x) = f(x) * (1 - f(x))$$

- Commonly used for models where we want to predict a probability as the function's output is between 0 and 1.
- Easy to cause vanishing gradient.

# Activation functions - Tanh

Tanh Function

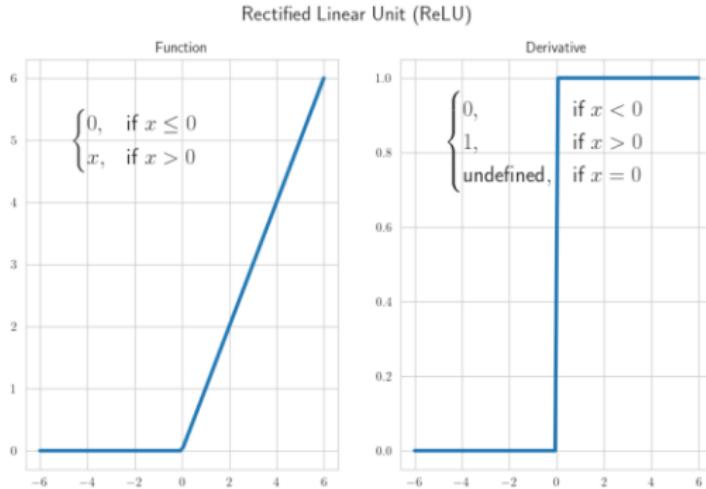


$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$f'(x) = 1 - f(x)^2$$

- The range of the tanh function is between -1 and 1.
- Good replacement of sigmoid to avoid vanishing gradient.

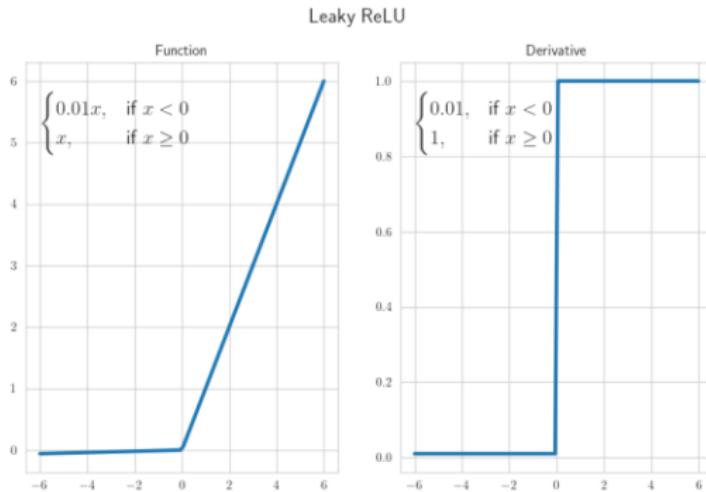
# Activation functions - ReLU



$$f(x) = \begin{cases} 0, & \text{if } x \leq 0 \\ x, & \text{if } x > 0 \end{cases}$$
$$f'(x) = \begin{cases} 0, & \text{if } x < 0 \\ 1, & \text{if } x > 0 \\ \text{undefined}, & \text{if } x = 0 \end{cases}$$

- Quick and simple calculation. The value range is between 0 and  $+\infty$ .
- At  $x = 0$ , the function is not differentiable because its left and right derivative are not equal.
- Negative values become zero which decreases the quality of training process.

# Activation functions – Leaky ReLU

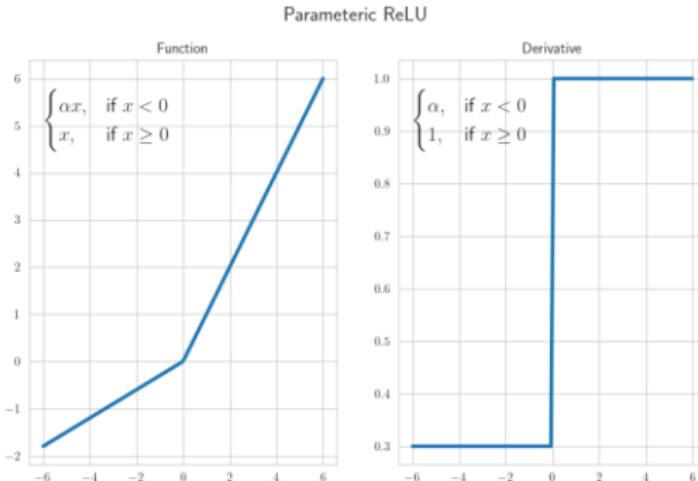


$$f(x) = \begin{cases} 0.01x, & \text{if } x \leq 0 \\ x, & \text{if } x > 0 \end{cases}$$

$$f'(x) = \begin{cases} 0.01, & \text{if } x < 0 \\ 1, & \text{if } x > 0 \\ \text{undefined}, & \text{if } x = 0 \end{cases}$$

- A variant of ReLU. Quick and simple calculation. The value range is between  $-\infty$  and  $+\infty$ . Usually, the parameter is often 0.01 or 0.1
- At  $x = 0$ , the function is not differentiable because its left and right derivative are not equal.

## Activation functions - Parametric ReLU

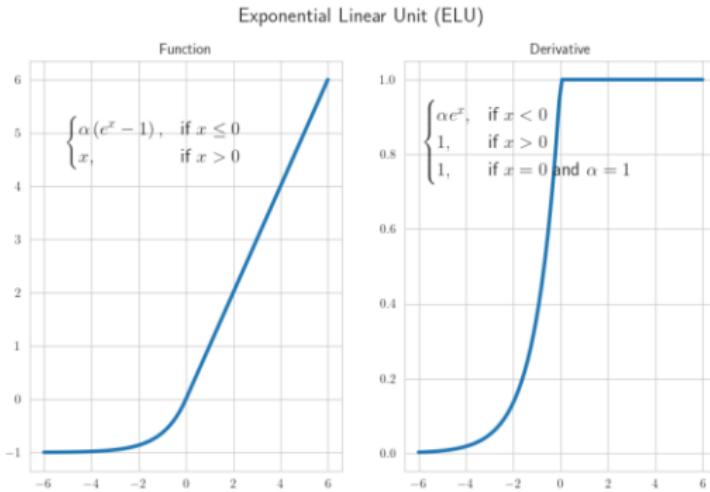


$$f(x) = \begin{cases} \alpha x, & \text{if } x \leq 0 \\ x, & \text{if } x > 0 \end{cases}$$

$$f'(x) = \begin{cases} \alpha, & \text{if } x < 0 \\ 1, & \text{if } x > 0 \\ \text{undefined}, & \text{if } x = 0 \end{cases}$$

- A variant of Leaky ReLU with parameter  $\alpha$ . The value range is between  $-\infty$  and  $+\infty$ .
- At  $x = 0$ , the function is not differentiable because its left and right derivative are not equal.

# Activation functions - ELU (Exponential Linear Unit)



$$f(x) = \begin{cases} \alpha(e^x - 1), & \text{if } x \leq 0 \\ x, & \text{if } x > 0 \end{cases}$$

$$f'(x) = \begin{cases} \alpha e^x, & \text{if } x < 0 \\ 1, & \text{if } x > 0 \\ 1, & \text{if } x = 0 \text{ and } \alpha = 1 \end{cases}$$

- Continuous and differentiable at all points, with given condition at  $x = 0$  and  $\alpha = 1$ .
- Using ELU often leads to a lower training time and a higher accuracy as compared to ReLU and its variants.

# Activation functions - Others

Built-in activation functions in Tensorflow:

`elu(...)`: Exponential Linear Unit.

`exponential(...)`: Exponential activation function.

`gelu(...)`: Applies the Gaussian error linear unit (GELU) activation function.

`get(...)`: Returns function.

`hard_sigmoid(...)`: Hard sigmoid activation function.

`linear(...)`: Linear activation function (pass-through).

`relu(...)`: Applies the rectified linear unit activation function.

`selu(...)`: Scaled Exponential Linear Unit (SELU).

`serialize(...)`: Returns the string identifier of an activation function.

`sigmoid(...)`: Sigmoid activation function,  $\text{sigmoid}(x) = 1 / (1 + \exp(-x))$ .

`softmax(...)`: Softmax converts a vector of values to a probability distribution.

`softplus(...)`: Softplus activation function,  $\text{softplus}(x) = \log(\exp(x) + 1)$ .

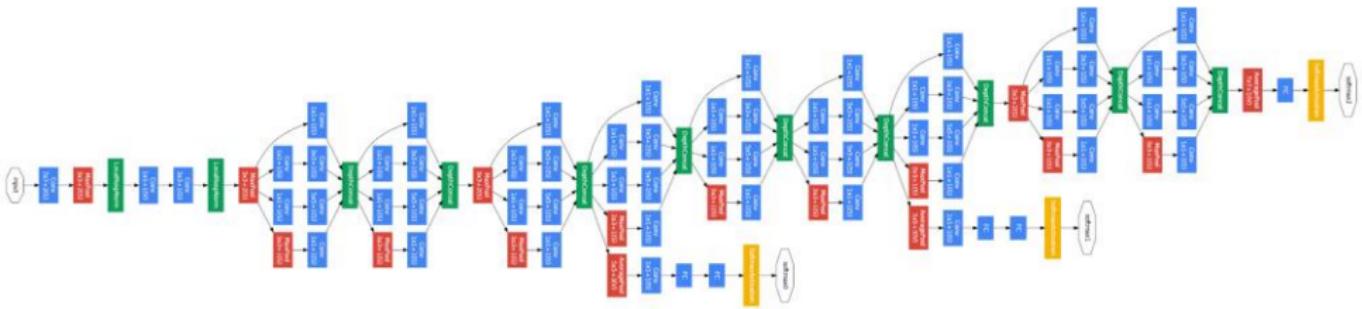
`softsign(...)`: Softsign activation function,  $\text{softsign}(x) = x / (\text{abs}(x) + 1)$ .

`swish(...)`: Swish activation function,  $\text{swish}(x) = x * \text{sigmoid}(x)$ .

`tanh(...)`: Hyperbolic tangent activation function.

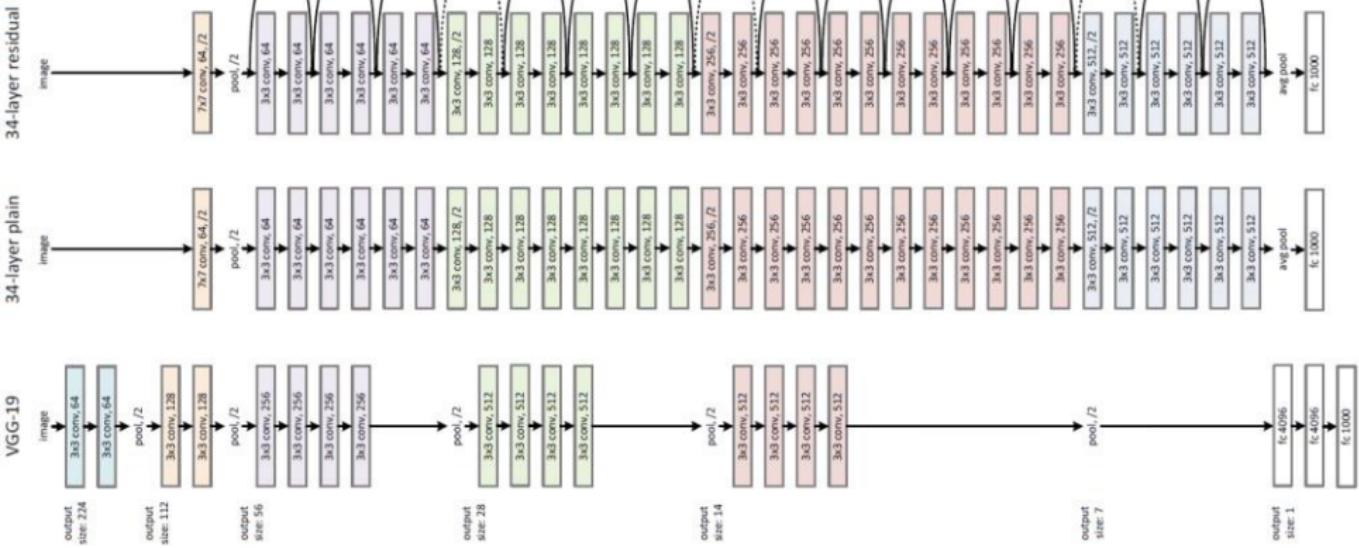
Reference: [https://www.tensorflow.org/api\\_docs/python/tf/keras/activations](https://www.tensorflow.org/api_docs/python/tf/keras/activations)

# Neural Network Architectures



GoogLe Net (Inception v1)

# Neural Network Architectures

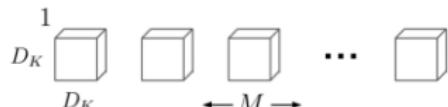


## ResNet

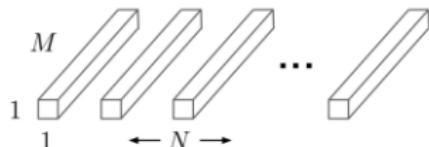
# Neural Network Architectures



(a) Standard Convolution Filters



(b) Depthwise Convolutional Filters



(c)  $1 \times 1$  Convolutional Filters called Pointwise Convolution in the context of Depthwise Separable Convolution

Type / Stride	Filter Shape	Input Size
Conv / s2	$3 \times 3 \times 3 \times 32$	$224 \times 224 \times 3$
Conv dw / s1	$3 \times 3 \times 32$ dw	$112 \times 112 \times 32$
Conv / s1	$1 \times 1 \times 32 \times 64$	$112 \times 112 \times 32$
Conv dw / s2	$3 \times 3 \times 64$ dw	$112 \times 112 \times 64$
Conv / s1	$1 \times 1 \times 64 \times 128$	$56 \times 56 \times 64$
Conv dw / s1	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 128$	$56 \times 56 \times 128$
Conv dw / s2	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 256$	$28 \times 28 \times 128$
Conv dw / s1	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 256$	$28 \times 28 \times 256$
Conv dw / s2	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 512$	$14 \times 14 \times 256$
$5 \times$ Conv dw / s1	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
	$1 \times 1 \times 512 \times 512$	$14 \times 14 \times 512$
Conv dw / s2	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
Conv / s1	$1 \times 1 \times 512 \times 1024$	$7 \times 7 \times 512$
Conv dw / s2	$3 \times 3 \times 1024$ dw	$7 \times 7 \times 1024$
Conv / s1	$1 \times 1 \times 1024 \times 1024$	$7 \times 7 \times 1024$
Avg Pool / s1	Pool $7 \times 7$	$7 \times 7 \times 1024$
FC / s1	$1024 \times 1000$	$1 \times 1 \times 1024$
Softmax / s1	Classifier	$1 \times 1 \times 1000$

## MobileNet v1

# Neural Network

What if building N-layer Neural Network  
with N is large?



## Summary

Fully connected layer.

Convolutional layer.

- 1x1 convolutional layer
- 1D, 2D, 3D convolutional layer.
- Transposed convolutional layer.
- Deconvolutional layer.

Normalization layer

- Batch normalization.
- Layer normalization.
- Instance normalization.
- Group normalization.

Pooling layer.

Dropout layer.

Activation functions.

Some simple Neural Network Architectures.

## Q&A

Thank you